
CS 3424 Systems Programming
Fall 2020

PRACTICE MIDTERM EXAM

[ANSWER KEY]

1. Answer true or false for the following statements:

- F The set of *extended* regular expressions includes the \$ anchor.
- F In Bash, \$# evaluates to the number of arguments passed into the script
- T Given the following bash command, the output printed will be “Yes”:
`COUNT=100; if [[$COUNT -gt 50]]; then echo Yes; else echo No; fi`
- F Given the following bash command, the output printed will be “No”:
`COUNT=100; if [$COUNT > 50]; then echo Yes; else echo No; fi`
- T Given the following bash command, the output printed will be “Yes”:
`COUNT=0; if [[COUNT > 50]]; then echo Yes; else echo No; fi`
- T sed uses basic regular expressions by default, and can use extended regular expressions by passing the -r or -E options
- F Bash will use the test utility to evaluate the conditional expression: `[[-r “$filename”]]`
- F The grep utility will search recursively through any specified directories
- T In AWK, \$4 contains the value of the fourth field in the current record
- F In Bash, an exit status of 1 represents success.
- T The extended regular expression {0,1} is functionally equivalent to the ? quantifier
- F Given `cmd1 && cmd2`, the second command will run regardless of the outcome of the first command
- F The regex pattern matching operator in AWK is “=~”
- F In Bash, \$0 contains a string representing the entire command line and all parameters

2. Write a Bash filename pattern (not a regular expression) that will only match the two files indicated in bold and underlined:

```
cs2123p1Driver.c  Makefile    p1abc123.o    p1OutExtra.txt
cs2123p1Driver.h  p1          p1Extra.txt  p1Out.log
cs2123p1Driver.o  p1abc123.c  p1Input.txt  p1Out.txt
```

```
p1[aEI]*.[to]*    -or-    p1[aEI]*.{o,txt}
```

3. Write a bash script that will take one or more directory names as arguments, and count the directories *under* each (that is, the number of subdirectories, subdirectories of subdirectories, and so on; do not count the arguments themselves).

Hint: you may use utilities we've discussed in class to help accomplish this.

```
#!/bin/bash

if [[ $# -eq 0 ]]; then
    echo "Usage: $0 <dir> [dir2 ... dirN]" 1>&2
    exit 1
fi

output=`find "$@" -type d | wc -l`
let count=$((output - $#))

echo "count = $count"
```

4. Write a sed command that will remove all lines longer than 80 characters.

```
/.\{81,\}/d
```

5. Write an AWK script that will perform the specified operations to the following data:

- Remove the last field from every record
- Compute and print the total count of all personnel, as well as the average age of the personnel

Example input:

```
Joe DiAlo          24 5023.00   180   890349013 07-19-1994
Christian Amun     44 8825.00   255   860028571 11-29-1974
Kale Martinez      32 5929.01   360   015993535 02-06-1986
John Doe           27 852.60    55    585921492 <no_dob> 00-00-0000
```

Example output:

Joe DiAlo 24 5023.00 180 890349013
 Christian Amunpour 44 8825.00 255 860028571
 Kale Martinez 32 5929.01 360 015993535
 John Doe 17 852.60 55 585921492 <no_dob>
 Average age = 29.25

```
1 {
2     $NF = ""
3     agesum += $3
4 }
5
6 END {
7     avgage = agesum / NR
8     print "Average age =", avgage
9 }
```

6. Write a sed command that will delete all leading blank lines at the top of a file.

```
./.,$!d
- alternative solution: -
0,././{h;d};H;z;x;s/^\\n//
```

7. Write a find command that will locate all regular files in your home directory and below that contain the phrase “versionxx” in their name, where each x represents a digit. For each file, run the “ls -la” command on them.

```
find ~ -name "*version[0-9][0-9]*" -type f -exec ls -la {} \;
```

8. Write a sed script that will:

- In all lines starting with “*PTY”, replace the string “*PTY” with “/*” characters and end the line with “*/”
- In the lines starting with “*PTY”, replace “X11R6” with “XFree86”
- Do not replace “X11R6” on any other lines.

```
/^\\*PTY/ {
s_/_\\*_
s_$_\\*_/_
s_X11R6_XFree86_g
}
```

9. Given an *excerpt* from a web log (named `part.log`), use a single pipeline to count the number of accesses made from the user with username root.

Hint: the command `wc -l` will count the number of lines it reads from standard input.

Time	User	Address	Req	Path
12740663209374.890	root	168.29.172.18	HTTP/GET	/www/img/*.gif
12740703137488.990	sam	129.82.45.103	HTTP/GET	/www/index.htm
12740897683420.004	hongyu	66.5.232.10	HTTP/POST	/www/root/profile
12741398476528.413	root	129.66.140.111	HTTP/GET	/home/cgi-bin/login.pl

```
grep " root " part.log | wc -l
grep -w root part.log | wc -l
awk '$2 == "root" { n++ } END { print n }' part.log
```

10. You run your own mail server on your home machine and you have had a number of sites attempting to relay mail through your server, dryad. So you want a simple sed script that will show the date of each relay attempt as well as the claimed source and destination accounts. **While there are many different types of logged information in this log file**, the lines relating to a rejected relay attempt is characterized by "Relay denied". These lines are similar to:

Feb 26 08:45:59 dryad Relay denied; from=<dlh@email.cq.cnt> to=gogo@ln.com.tw

The corresponding output should be:

Feb 26 08:45:59 dlh@email.cq.cnt -> gogo@ln.com.tw

Write a sed script that will transform each line of input to produce the specified output.

```
/Relay denied/ {
    s/dryad Relay denied; from=<>//
    s/> to=/ -> /
}
-or-
/Relay denied/ s/> to=/ -> /
/Relay denied/ s/dryad Relay denied; from=<//
```

11. Will the following grep command find a match, no match, or produce an error?

```
echo xyz | grep -E "(dog|cat|bat)*y(horse|pig)?"
```

Match: xyz

12. Write a bash script that takes a filename as a parameter. After verifying that the file exists and is readable, use the “wc -l” command in order to obtain a line count for each file. Find the longest file (in terms of their number of lines), and print this file’s name, as well as its length in lines. (*Don’t worry about breaking ties for files with equal numbers of lines*).

```
#!/bin/bash

if [ $# -lt 1 ]; then
    echo "Usage: $0 <file1> [file2] ... [fileN]" 2>&1
    exit 1
fi

let maxlines=0
for file; do

    # Old-style conditional expression first, new-style second
    #if [ \( ! -r "$file" \) -o \( -d "$file" \) ]; then continue; fi
    if [[ ( ! -r "$file" ) || ( -d "$file" ) ]]; then continue; fi

    length=$(wc -l < "$file")

    if [ $length -gt $maxlines ]; then
        maxlines=$length
        maxfile="$file"
    fi
done

echo "longest file:  $maxfile"
echo "length (lines): $maxlines"
```

13. Write a sed command that will prepend the following line prior to any line that might contain a social-security number. I.e., any line that matches this pattern:

[0-9]{3}-[0-9]{2}-[0-9]{4}:

*** the following line requires redaction ***

Example input:

...the patient’s SSN is 435-32-7992, DOB is 7/12/89, ...

Example output:

*** the following line requires redaction ***

...the patient’s SSN is 435-32-7992, DOB is 7/12/89, ...

/[0-9]{3}-[0-9]{2}-[0-9]{4}/ \

*** line requires redaction ***

-or-

/[0-9]{3}-[0-9]{2}-[0-9]{4}/ i *** line requires redaction ***

-or-

`/[0-9]{3}-[0-9]{2}-[0-9]{4}/ s/^*** line requires redaction ***\n/`

14. Write a sed script that will print the real names of up to the first 100 users whose shell is `bash` (the real name is the fifth colon-delimited field; their shell is specified in the last field).

Example input file *excerpt* (pretend this file is actually hundreds of lines long):

```
huunguyen:x:1021:1020:Huu Nguyen:/home/huunguyen:/usr/bin/nologin
tianyiliu:x:1022:1022:Tianyi Liu:/home/tianyiliu:/bin/tcsh
quadewarren:x:991:1023:Kyle Landry:/home/quadewarren:/bin/bash
yuyutang:x:1024:1024:Yuyu Tang:/home/yuyutang:/usr/bin/nologin
xuelingzhang:x:725:1025:Xueling Zhang:/home/xuelingzhang:/bin/bash
dalton:x:1026:1026:TJ Bashwami:/home/bashwami:/bin/tcsh
michael:x:980:1027:Michael Thompson:/home/michael:/usr/bin/nologin
```

...

Example output:

```
Huu Nguyen
Tianyi Liu
Kyle Landry
Yuyu Tang
Xueling Zhang
TJ Bashwami
Michael Thompson
```

...

`1,100 s/.*:.*:.*:.*:.*:(.*):.*:.*$/\1/p # use only with -n`

-or-

```
1,100 {
    s/.*:(.*):.*:.*$/\1/
    p
}
d
```

-or-

```
s/.*:.*:.*:.*:.*:(.*):.*:.*$/\1/
101,$ d
```

```
1,100 {
    s/.*:(.*):.*/\1/
    p
}
```

```
1 BEGIN { FS = ":" }
2 1,100 {
3     if($7 == "/bin/bash") {
4         print $5
5     }
```

b. Write a sed script that will print the real names of up to the first 100 users **whose shell is bash** (the real name is the fifth colon-delimited field; their shell is specified in the last field).
Example input file *excerpt* (pretend this file is actually hundreds of lines long):

```
...
1,100 {
    s/.*:.*:.*:.*:.*:(.*)\:.*:\/bin\/bash$/\1/p # use only with -n
}
```

```
1,100 {
    /:/bin\bash$/ s/.*:.*:.*:.*:(.*):.*:.*/\1/p # use only with -n
}
d
```

-or-

```
1,100 {  
    /:\bin\bash$/ {  
        s/.*:.*):.*:\bin\bash$/\1/  
        p  
    }  
}  
d
```

-or-

```
:\bin\bash$/! d  
:\bin\bash$/ s/.*:.*):.*:\bin\bash$/\1/  
101,$ d
```

17. Find the match for **group 1** in the given string and regular expression. If there is no match, write “NO MATCH.” Assume the use of extended regular expressions.

String	Pattern	Group 1
abcd123+456efgh	/([a-c]+[d-q].)/	abcd1
abcd123+456efgh	/[0-9]*([a-d]+)/	abcd
abcd123+456efgh	/(^[0-9]*.{3}).*/	abc
abcd123+456efgh	/d.([0-9]*)([1-3])/	2
abcd123+456efgh	/([e-g678]+).\$/	6efg